

```
! pip install torchvision
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheelhouse/pypi
Requirement already satisfied: torchvision in /usr/local/lib/python3.7/dist-packages (0.11.2)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.7/dist-packages (9.0.1)
Requirement already satisfied: torch==1.12.1 in /usr/local/lib/python3.7/dist-packages (1.12.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (1.24.2)
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (2.31.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (4.5.0)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (1.25.11)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (2022.9.24)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (3.7.4)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (3.4)
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
import math
import copy
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision.models as models
import torchvision.transforms as transforms
import torch.optim as optim
```

```
from PIL import Image
```

```
import glob
import matplotlib.image as mpimg
```

```
df = pd.read_csv('/content/drive/MyDrive/artists.csv')
```

```
df.head()
```

```

    id      name      years      genre      nationality      name
0  0      Amedeo      1884 -      Expressionism      Italian      Amedeo
    0      Modigliani      1920
1  1      Vasiliy      1866 -      Expressionism,Abstractionism      Russian      Wassily Was
    1      Kandinskiy      1944
2  2      Diego Rivera      1886 -      Social Realism,Muralism      Mexican      Diego
    2      1957
df.set_index('id', inplace=True)
3  3      Claude Monet      1926      Impressionism      French
df.head()
```

	name	years	genre	nationality	
id					
0	Amedeo Modigliani	1884 - 1920	Expressionism	Italian	Amedeo
1	Vasiliy Kandinskiy	1866 - 1944	Expressionism,Abstractionism	Russian	Wassily Was
2	Diego Rivera	1886 - 1957	Social Realism,Muralism	Mexican	Diego
3	Claude Monet	1840 -	Impressionism	French	Oscar-Claude M

```

model= models.vgg19(pretrained = True).features

/usr/local/lib/python3.7/dist-packages/torchvision/models/_utils.py:209: UserWarning:
f"The parameter '{pretrained_param}' is deprecated since 0.13 and will be removed in a future
/usr/local/lib/python3.7/dist-packages/torchvision/models/_utils.py:223: UserWarning:
warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/vgg19-dcbb9e9d.pth" to /root/.cache/torch/hub
100%          548M/548M [00:31<00:00, 29.5MB/s]

print(model)

Sequential(
  (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU(inplace=True)
  (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU(inplace=True)
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (6): ReLU(inplace=True)
  (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (8): ReLU(inplace=True)
  (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)
```

```

(11): ReLU(inplace=True)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace=True)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace=True)
(16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(17): ReLU(inplace=True)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace=True)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): ReLU(inplace=True)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(26): ReLU(inplace=True)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): ReLU(inplace=True)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(33): ReLU(inplace=True)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(35): ReLU(inplace=True)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)

```

```

for param in model.parameters():
    param.requires_grad_(False)

```

```

dev = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(dev)

```

```

Sequential(
  (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU(inplace=True)
  (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU(inplace=True)
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (6): ReLU(inplace=True)
  (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (8): ReLU(inplace=True)
  (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (11): ReLU(inplace=True)
  (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (13): ReLU(inplace=True)
  (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (15): ReLU(inplace=True)
  (16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)

```

```

(17): ReLU(inplace=True)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace=True)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): ReLU(inplace=True)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(26): ReLU(inplace=True)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): ReLU(inplace=True)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(33): ReLU(inplace=True)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(35): ReLU(inplace=True)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
)

```

```

from torchvision.ops.poolers import torchvision
def load_img(path, max_size = 400, shape = None):
    image = Image.open(path).convert('RGB')

    if max(image.size) > max_size:
        size = max_size
    else:
        size = max(image.size)

    if shape is not None:
        size = shape

```

```

in_trans = transforms.Compose([transforms.Resize(size),
                               transforms.ToTensor(),
                               transforms.Normalize((0.485, 0.456, 0.406),
                                                    (0.229, 0.224, 0.225))])

image = in_trans(image)[:3,:,:].unsqueeze(0)

return image

```

```

from PIL import Image
import glob
image_lst = []
for filename in glob.glob('/content/drive/MyDrive/Dataset/*.jpg'):

```

```
im = Image.open(filename)
image_lst.append(im)

data = load_img('/content/drive/MyDrive/Dataset/Albrecht_Dürer_102.jpg').to(dev)
style_data = load_img('/content/drive/MyDrive/Dataset/Amedeo_Modigliani_109.jpg').to(dev)


def convert(tensor):
    image = tensor.to("cpu").clone().detach()
    image = image.numpy().squeeze()
    image = image.transpose(1,2,0)
    image = image * np.array((0.229, 0.224, 0.225)) + np.array((0.485, 0.456, 0.406))
    image = image.clip(0, 1)

    return image

fig, (ax1, ax2) = plt.subplots(1, 2, figsize = (20, 10))

ax1.imshow(convert(data))
ax1.set_title("Content Image", fontsize = 20)
ax2.imshow(convert(style_data))
ax2.set_title("Style Image", fontsize = 20)
plt.show()
```

Content Image



## ▼ Model



```
def features(image, model, layers= None):
    if layers is None:
        layers = {'0': 'conv1_1',
                  '5': 'conv2_1',
                  '10': 'conv3_1',
                  '19': 'conv4_1',
                  '21': 'conv4_2',
                  '28': 'conv5_1'}

    features = {}
    x= image
    for name, layer in model._modules.items():
        x = layer(x)
        if name in layers:
            features[layers[name]] = x

    return features
```



```
# gram matrix
def gram_matrix(tensor):
    _, d, h, w = tensor.size()

    tensor = tensor.view(d, h*w)
    gram = torch.mm(tensor, tensor.t())
    return gram

cont_features = features(data, model)
style_feat = features(style_data, model)

style_grams = {layer: gram_matrix(style_feat[layer]) for layer in style_feat}
target = data.clone().requires_grad_(True).to(dev)

# Loss and Weights

style_weights = {'conv1_1': 1.,
                  'conv2_1': 0.75,
                  'conv3_1': 0.2,
                  'conv4_1': 0.2,
                  'conv5_1': 0.2}
```

```
data_weights= 1
style_weight = 1e9

#Total loss

show_every = 400

optimizer = optim.Adam([target], lr = 0.003)
steps = 2000

for ii in range(1, steps+1):
    target_feat = features(target, model)
    cont_loss = torch.mean((target_feat['conv4_2'] - cont_features['conv4_2'])**2)

    style_loss = 0
    for layer in style_weights:
        target_feats = target_feat[layer]
        target_gram = gram_matrix(target_feats)
        _, d, h, w = target_feats.shape

        style_gram = style_grams[layer]
        layer_style_loss = style_weights[layer]* torch.mean((target_gram - style_gram)**2)
        style_loss += layer_style_loss / (d * h * w)

    total_loss = data_weights * cont_loss + style_weight * style_loss

    optimizer.zero_grad()
    total_loss.backward()
    optimizer.step()

    if ii % show_every == 0:
        print('Total Loss: ', total_loss.item())
        plt.imshow(convert(target))
        plt.axis('off')
        plt.grid(False)
        plt.figure(figsize = (7, 7))
        plt.show()
```

Total Loss: 91320958976.0



<Figure size 504x504 with 0 Axes>

Total Loss: 63734923264.0



<Figure size 504x504 with 0 Axes>

Total Loss: 39968530432.0



<Figure size 504x504 with 0 Axes>

Total Loss: 24250875904.0







<Figure size 504x504 with 0 Axes>  
Total Loss: 14284768256.0

```
fig, (ax1, ax2, ax3) = plt.subplots(1, 3, figsize=(15, 15))
ax1.imshow(convert(data))
ax1.set_title("Content Image", fontsize = 20)
ax2.imshow(convert(style_data))
ax2.set_title("Style Image", fontsize = 20)
ax3.imshow(convert(target))
ax3.set_title("Stylized Target Image", fontsize = 20)
ax1.grid(False)
ax2.grid(False)
ax3.grid(False)

ax1.set_xticks([])
ax1.set_yticks([])
ax2.set_xticks([])
ax2.set_yticks([])
ax3.set_xticks([])
ax3.set_yticks([])

plt.show()
```



Content Image



Style Image



Stylized Target Image



[Colab paid products](#) - [Cancel contracts here](#)

